

# Predicting Banking Loan Defaults

Team Members: Dheekshith, Ninad, Mouneesh, Mahith

Github Link:

[https://github.com/f20212948/Gramener\\_Training/tree/main/Project1\\_Banking\\_Default\\_Group1](https://github.com/f20212948/Gramener_Training/tree/main/Project1_Banking_Default_Group1)

Kaggle Data Link:

[https://www.kaggle.com/datasets/yasserh/loan-default-dataset/data?select=Loan\\_Default.csv](https://www.kaggle.com/datasets/yasserh/loan-default-dataset/data?select=Loan_Default.csv)

This project focuses on building and comparing machine learning models to predict loan defaults (the **Status** column) using the provided banking loan dataset. The process involves comprehensive data analysis, preprocessing (handling missing values, skewness, and categorical features), feature selection, managing class imbalance, and training various classification models, with a focus on **XGBoost**.

## Project Structure and Flow

The analysis follows a standard machine learning workflow:

1. **Data Loading and Initial Exploration:** Load the dataset and perform initial checks for shape, data types, and missing values.
2. **Exploratory Data Analysis (EDA):** Visualize distributions, check for outliers, analyze correlations, and examine the relationship between features and the target variable (**Status**).
3. **Data Preprocessing and Transformation:**
  - **Handling Missing Values:** Missing values in numerical columns are imputed using the **median strategy** (SimpleImputer).
  - **Skewness Transformation:** Highly skewed numerical features are transformed using  $\log(1+x)$  to make their distributions more Gaussian-like.
  - **Categorical Encoding:** Object type columns are imputed with the **most frequent value** and then transformed using **One-Hot Encoding**.
4. **Feature Selection and Scaling:**
  - A smaller sample of the data is used for **Hyperparameter Tuning** to save time.
  - Numerical features are **Standard Scaled**.
  - The **XGBoost Classifier** is used on the sampled data to identify the **Top 30 Most Important Features**.
5. **Handling Class Imbalance:** The full training data is subjected to **Random Under-Sampling (RUS)** to balance the majority (non-default) and minority (default) classes.
6. **Model Training and Evaluation:** Three main models are trained, tuned, and evaluated: **Logistic Regression**, **Random Forest Classifier**, and **XGBoost Classifier**.
  - **Hyperparameter Tuning** is performed using **RandomizedSearchCV** on the sampled data to find optimal parameters (e.g., **C** for Logistic Regression, **max\_depth** for Random Forest and XGBoost).

- Final models are trained on the **under-sampled training data** and evaluated on a separate **test set** and **validation set**.
7. **Conclusion:** Compare model performance using metrics like **ROC AUC**, **Precision-Recall AUC (PR AUC)**, and **Accuracy**.
- 

## Key Code Logic Details

Section	Technique Used	Purpose
<b>Missing Values</b>	<code>SimpleImputer(strategy='median')</code>	To fill missing values in numerical columns, as the median is less sensitive to outliers.
<b>Data Cleaning</b>	Custom <code>type_transform</code> and <code>purpose_transform</code> functions	To convert cryptic categorical codes (e.g., 'type1', 'p1') into descriptive strings (e.g., 'Conventional Loans', 'Home Purchase') for better interpretability.
<b>Skewness</b>	<code>numpy.log1p(x)</code>	Applies a $\log(1+x)$ transformation to reduce positive skewness in numerical features.
<b>Categorical Processing</b>	<code>ColumnTransformer</code> with <code>SimpleImputer('most_frequent')</code> and <code>OneHotEncoder</code>	Handles both imputation for categorical features and converting them into a format suitable for machine learning models.

<b>Hyperparameter Tuning</b>	RandomizedSearchCV	Efficiently searches a wide range of hyperparameters for <b>Logistic Regression</b> , <b>Random Forest</b> , and <b>XGBoost</b> to find the best configuration based on <b>ROC AUC</b> score.
<b>Feature Selection</b>	XGBClassifier.feature_importances_	Uses the tuned XGBoost model's feature importance to select the <b>Top 30 most impactful features</b> , reducing dimensionality for final models.
<b>Class Imbalance</b>	RandomUnderSampler from <code>imblearn</code>	Reduces the number of samples in the majority class to balance the dataset, which is crucial for improving the prediction of the minority class (loan defaults).
<b>Final Evaluation</b>	ROC AUC, PR AUC, Classification Report	These metrics are critical for imbalanced classification tasks. <b>PR AUC</b> is especially important as it focuses on the performance on the positive class (default).

## How to Run the Code

- Save the Code:** Ensure the Python script containing the code is saved (e.g., `loan_prediction.py`).
- Acquire Data:** Ensure the dataset `Loan_Default.csv` is in the same directory as the script.

**Execute:** Run the script from your terminal:

Bash

```
python loan_prediction.py
```

3. The script will output the following:
  - Multiple EDA plots (histograms, boxplots, heatmaps).
  - Best hyperparameters and best ROC AUC scores from [RandomizedSearchCV](#).
  - Feature importance list.
  - Classification Reports and performance scores (Accuracy, ROC AUC, PR AUC) for each model on the test and validation sets.
  - The final ROC Curve and Precision-Recall Curve plots comparing all three models.

## Required Installations

This project requires several common Python libraries for data manipulation, machine learning, and visualization.

You can install all necessary packages using [pip](#):

Bash

```
pip install pandas numpy matplotlib seaborn scikit-learn xgboost imblearn
```

Library	Purpose
<b>pandas</b>	Data manipulation and analysis.
<b>numpy</b>	Numerical operations, especially for mathematical transformations.
<b>matplotlib</b>	Creating static, interactive, and animated visualizations.
<b>seaborn</b>	High-level interface for drawing attractive statistical graphics.
<b>scikit-learn</b>	Machine learning framework (models, preprocessing, evaluation).
<b>xgboost</b>	Highly efficient Gradient Boosting implementation.
<b>imblearn</b>	Handling imbalanced datasets (specifically, <a href="#">RandomUnderSampler</a> ).

---

## Summary of Model Performance

Based on the test set evaluation after addressing class imbalance and feature selection, the **XGBoost Classifier** is the recommended model:

Model	Accuracy	ROC AUC	PR AUC
<b>Logistic Regression</b>	0.867	0.831	0.757
<b>Random Forest</b>	0.855	0.806	0.735
<b>XGBoost Classifier (Best)</b>	<b>0.869</b>	<b>0.842</b>	<b>0.768</b>

The XGBoost model demonstrates the best overall ability to distinguish between loan defaulters and non-defaulters.